
Key-Value Repositories

Knowledge Objectives

1. Explain the need for key-value stores
2. Elaborate on the four goals of key-value stores
3. Define what is a schemaless database
4. Describe how key-value stores improve performance by means of parallelism

Understanding Objectives

1. Explain the two main consequences of schemaless databases

Application Objectives

1. Model simple schemaless databases

AN EXAMPLE OF KEY-VALUE ARCHITECTURE

HADOOP, BIGTABLE AND MAPREDUCE

Key-Values: A Piece of History


- ❑ Key-values were born as a desperate answer to the RDBMS limitations
- ❑ It is widely assumed that Google is the father of Key-value stores
 - ❑ Hadoop File System
 - ❑ *The Google File System* (2003)
 - ❑ MapReduce
 - ❑ *Simplified Data Processing on Large Clusters* (2004)
 - ❑ HBase
 - ❑ *A Distributed Storage System for Structured Data* (2006)

Google Ecosystem

- ❑ High-performance is mainly achieved by means of parallelism
 - Divide-and-conquer principle
- ❑ MapReduce
 - It is a query language that provides parallelism in a transparent manner




Query Language



Database

Google Ecosystem

- ❑ High-performance is mainly achieved by means of parallelism
 - Divide-and-conquer principle
- ❑ MapReduce
 - It is a query language that provides parallelism in a transparent manner



A DISTRIBUTED FILE SYSTEM

THE FILE SYSTEM: HADOOP

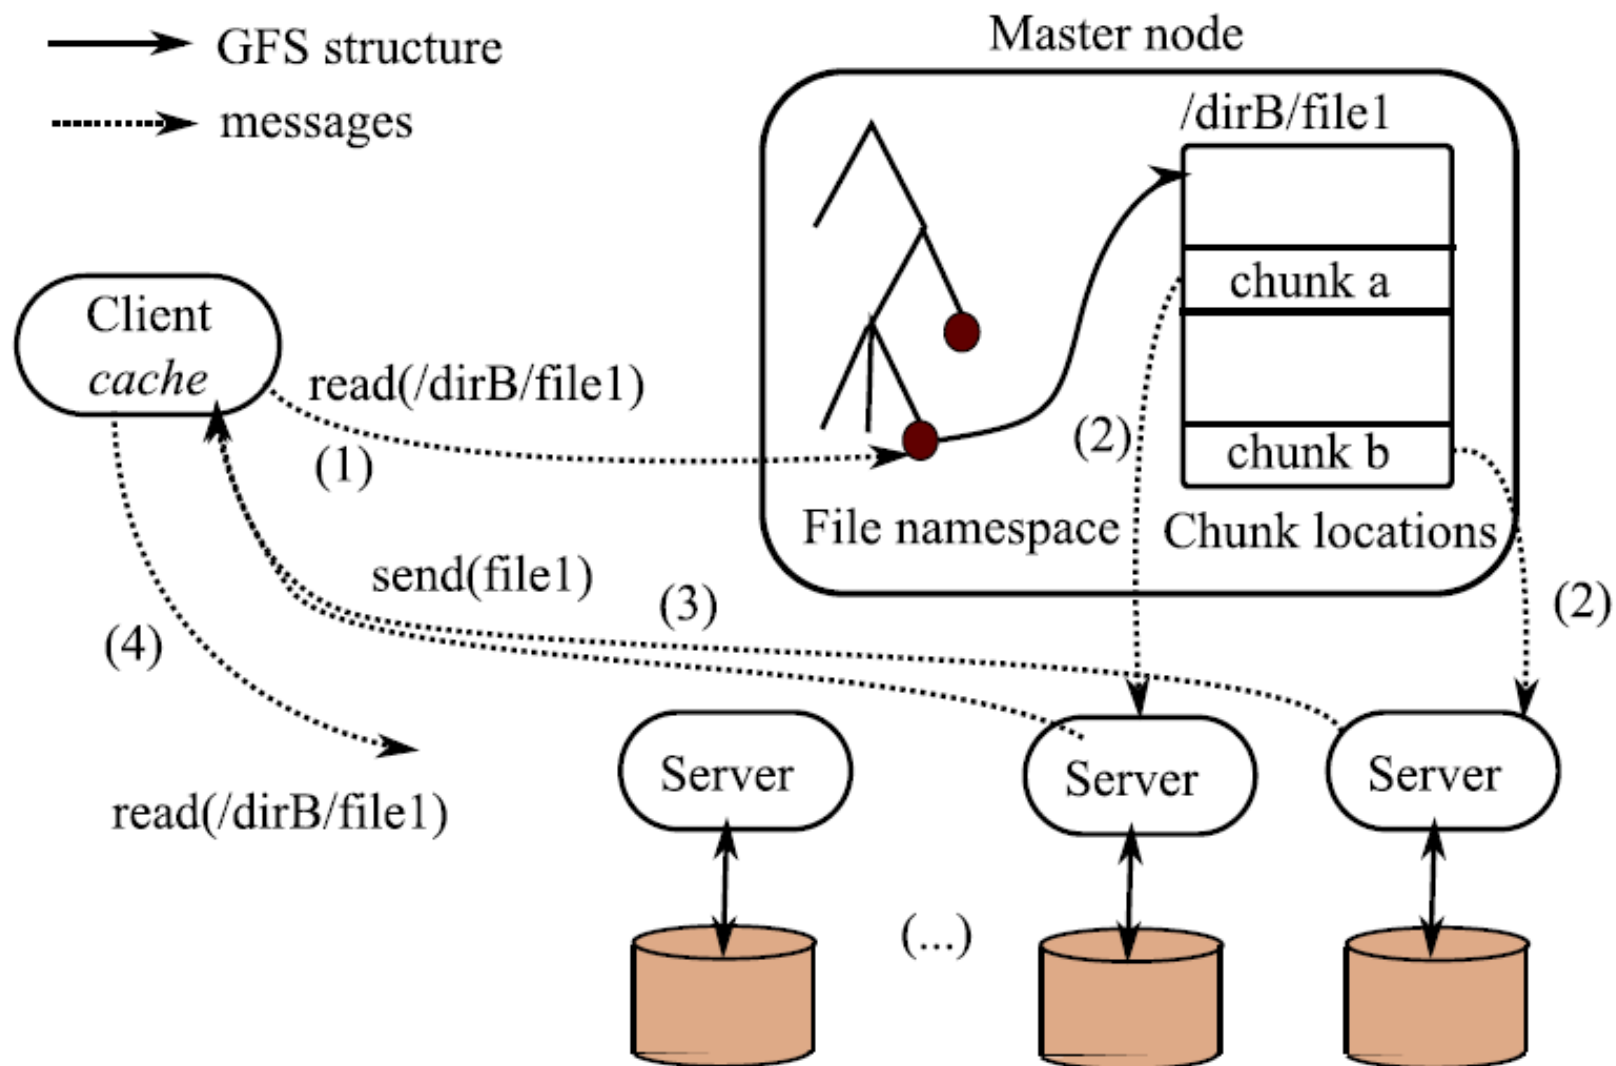
Hadoop File System (HDFS)

- ❑ Apache Hadoop (<http://hadoop.apache.org>)
 - Based on Google File System (GFS)
- ❑ Designed to meet the following requirements
 - Handle very large collections of unstructured to semi-structured data
 - Data collections are written once and read many times
 - The infrastructure underlying consists of thousands of connected machines with high failure probability
- ❑ Traditional network file systems do partially fulfil these requirements
 - Operating Systems Vs. Database Management System
 - ❑ Balancing *query* load (e.g., by means of *fragmentation and replication*) boosts availability and reliability
 - HDFS: Equal-sized file chunks evenly distributed

HDFS in a Nutshell (I)

- ❑ A single master (coordinator)
 - Receives client connections
 - Maintains the description of the global file system namespace
 - Keeps track of file chunks (default: 128Mb)
- ❑ Many servers
 - Receive file chunks and store them
- ❑ A single master design forfeits availability and scalability
 - Availability and reliability: Recovery system
 - ❑ Replication (a chunk always in 3 servers, by default)
 - ❑ Monitors the system with heartbeat messages to detect failures as soon as possible
 - ❑ Specific recovery system to protect the master
 - Scalability: Client cache

HDFS in a Nutshell (II)



HDFS File Formats

□ Parquet

- Native columnar storage
- Rich header metadata (including schema information)
- Supports advanced block compression techniques
- Recognised by most popular processing engines, such as Spark
 - It allows SQL-like querying (e.g., Spark SQL)
 - Selections and Projections can be pushed down to disk
 - Statistics used to skip whole fragments

□ Avro

- Native row-oriented storage
- Rich header metadata (including schema information)
- Supports advanced block compression techniques
- Recognised by most popular processing engines, such as Spark
 - It allows SQL-like querying (e.g., Spark SQL)

□ Text / CSV / JSON formats

- Do not support block compression
- Fix-sized splitting with no metadata
- Not advisable for in-Hadoop processing
- Typically used to store *raw* data

□ ... and many others: Arrow, ORC Files, Sequence Files, etc.

AN EXAMPLE OF KEY-VALUE ARCHITECTURE

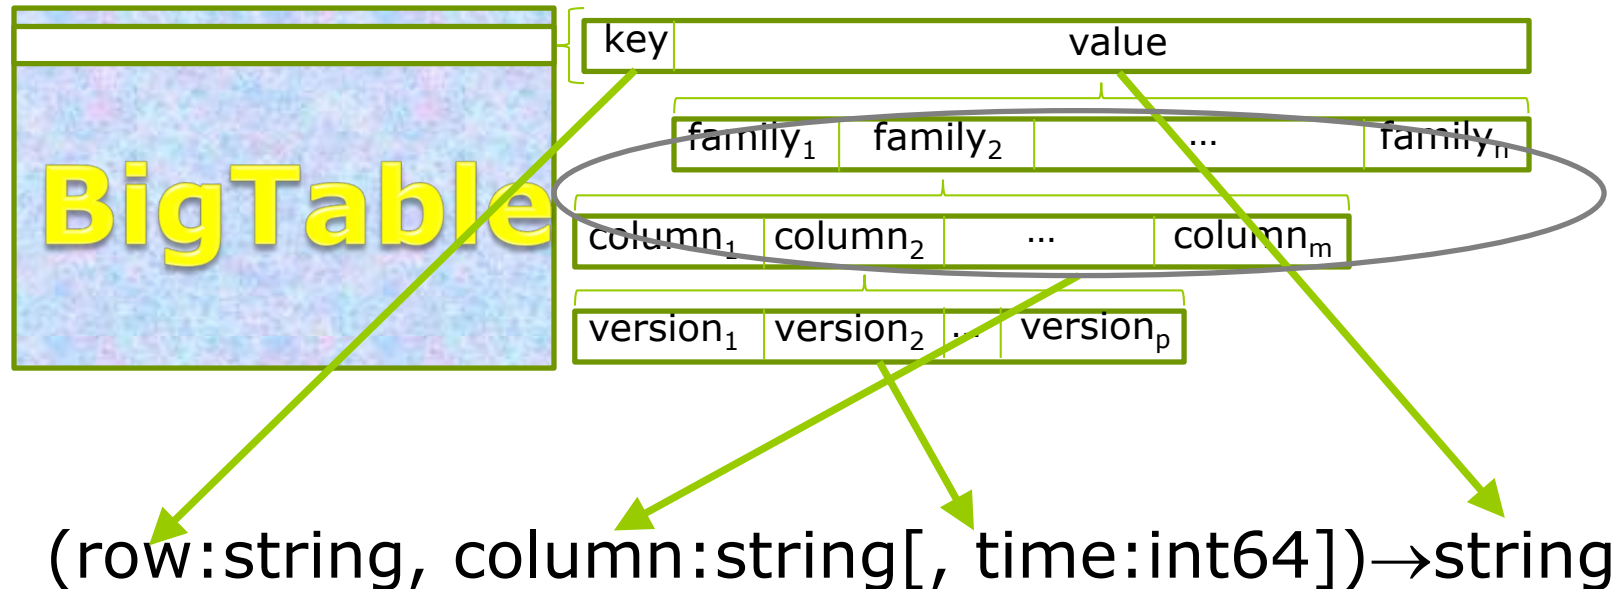
THE DATABASE: HBASE

HBase

- ❑ Apache HBase (<http://hbase.apache.org>)
 - Based on Google's BigTable
- ❑ Designed to meet the following requirements
 - Access specific data out of petabytes
 - It must support key search, range search and high throughput for file scans
 - It must support single row transactions
- ❑ *Do it yourself* database... own decisions regarding:
 - Data structure
 - Concurrency
 - Recovery
 - CAP trade-off
 - Etc.
- ❑ In short, it is a **distributed index cluster** on top of HDFS
 - Distributed B+
 - Tuples are lexicographically sorted according to the key

Schema Elements

- ❑ Stores tables (collections) and rows (instances)
 - Data is indexed using row and column names (arbitrary strings)
- ❑ Treats data as **uninterpreted** strings (without data types)
- ❑ Each cell of a BigTable can contain multiple versions of the same data
 - Stores different versions of the same values in the rows
 - Each version is identified by a timestamp
 - ❑ Timestamps can be explicitly or automatically assigned



Activity: Key-Value Design

❑ **Objective:** Learn the basic design principles of key-value stores

❑ **Tasks:**

1. (20') By pairs, solve the following exercise

- Model in HBase

the lineitem and order tables

- Model the whole schema

2. (5') Discussion

Q1

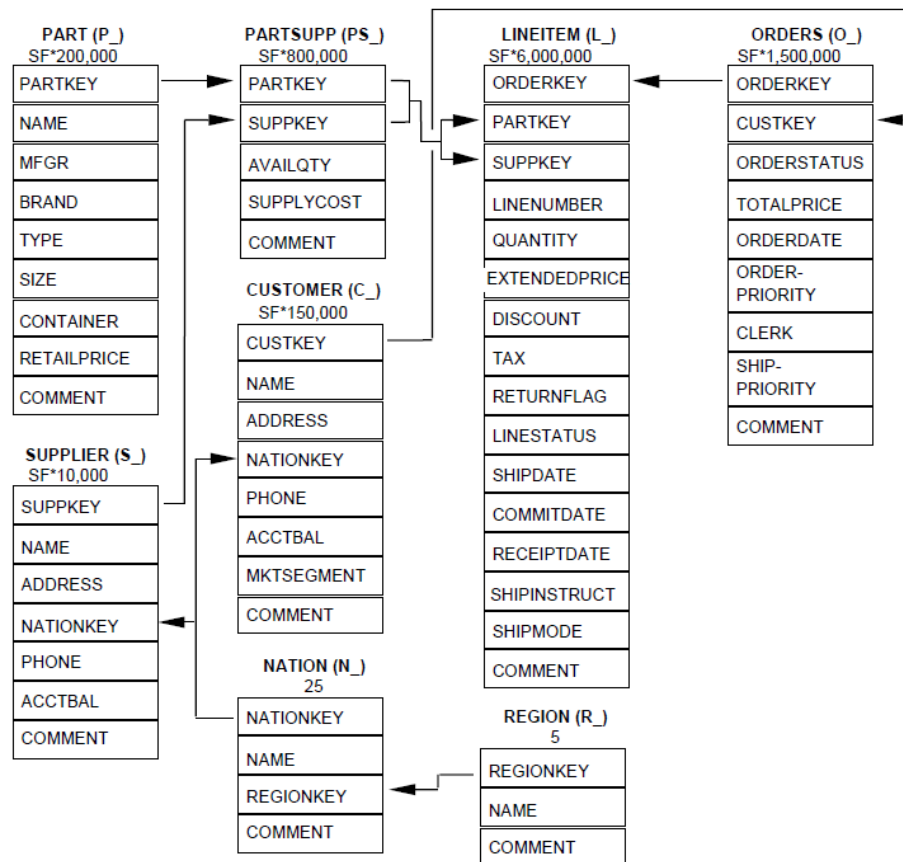
```
SELECT l_orderkey, sum(l_extendedprice*(1-
l_discount)) as revenue, o_orderdate,
o_shippriority
```

```
FROM customer, orders, lineitem
```

```
WHERE c_mktsegment = '[SEGMENT]' AND c_custkey =
o_custkey AND l_orderkey = o_orderkey AND
o_orderdate < '[DATE]' AND l_shipdate > '[DATE]'
```

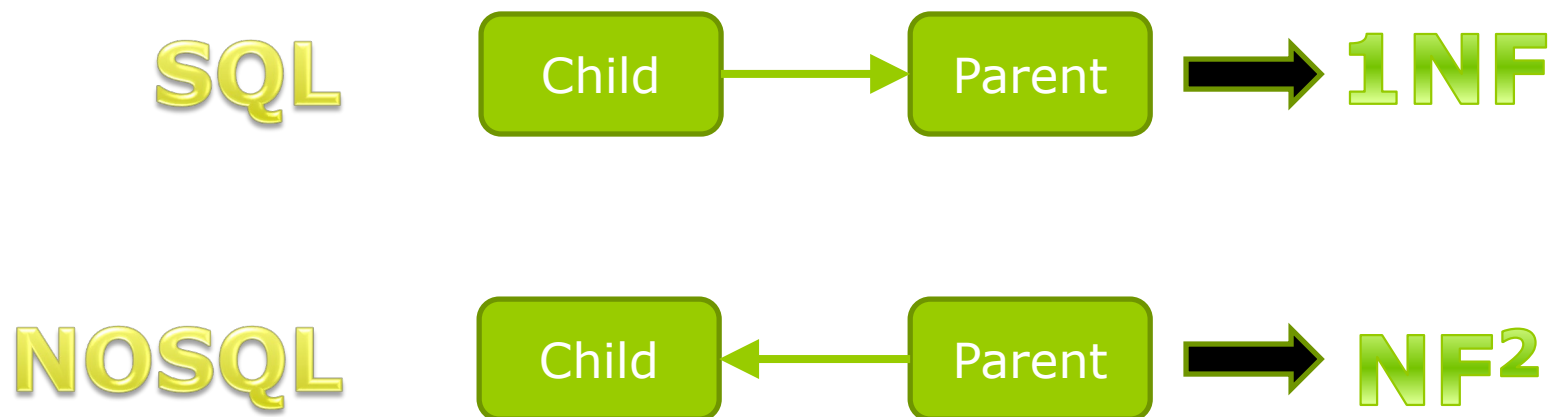
```
GROUP BY l_orderkey, o_orderdate, o_shippriority
```

```
ORDER BY revenue desc, o_orderdate;
```



TPC-H Benchmark

Just Another Point of View



Just Another Point of View




Just Another Point of View



HBase Shell

- ❑ ALTER <tablename>, <columnfamilyparam>
- ❑ COUNT <tablename>
- ❑ CREATE TABLE <tablename>
- ❑ DESCRIBE <tablename>
- ❑ DELETE <tablename>, <rowkey>[, <columns>]
- ❑ DISABLE <tablename>
- ❑ DROP <tablename>
- ❑ ENABLE <tablename>
- ❑ EXIT
- ❑ EXISTS <tablename>
- ❑ GET <tablename>, <rowkey>[, <columns>]
- ❑ LIST
- ❑ PUT <tablename>, <rowkey>, <columnid>, <value>[, <timestamp>]
- ❑ SCAN <tablename>[, <columns>]
- ❑ STATUS [{summary|simple|detailed}]
- ❑ SHUTDOWN

Physical Implementation

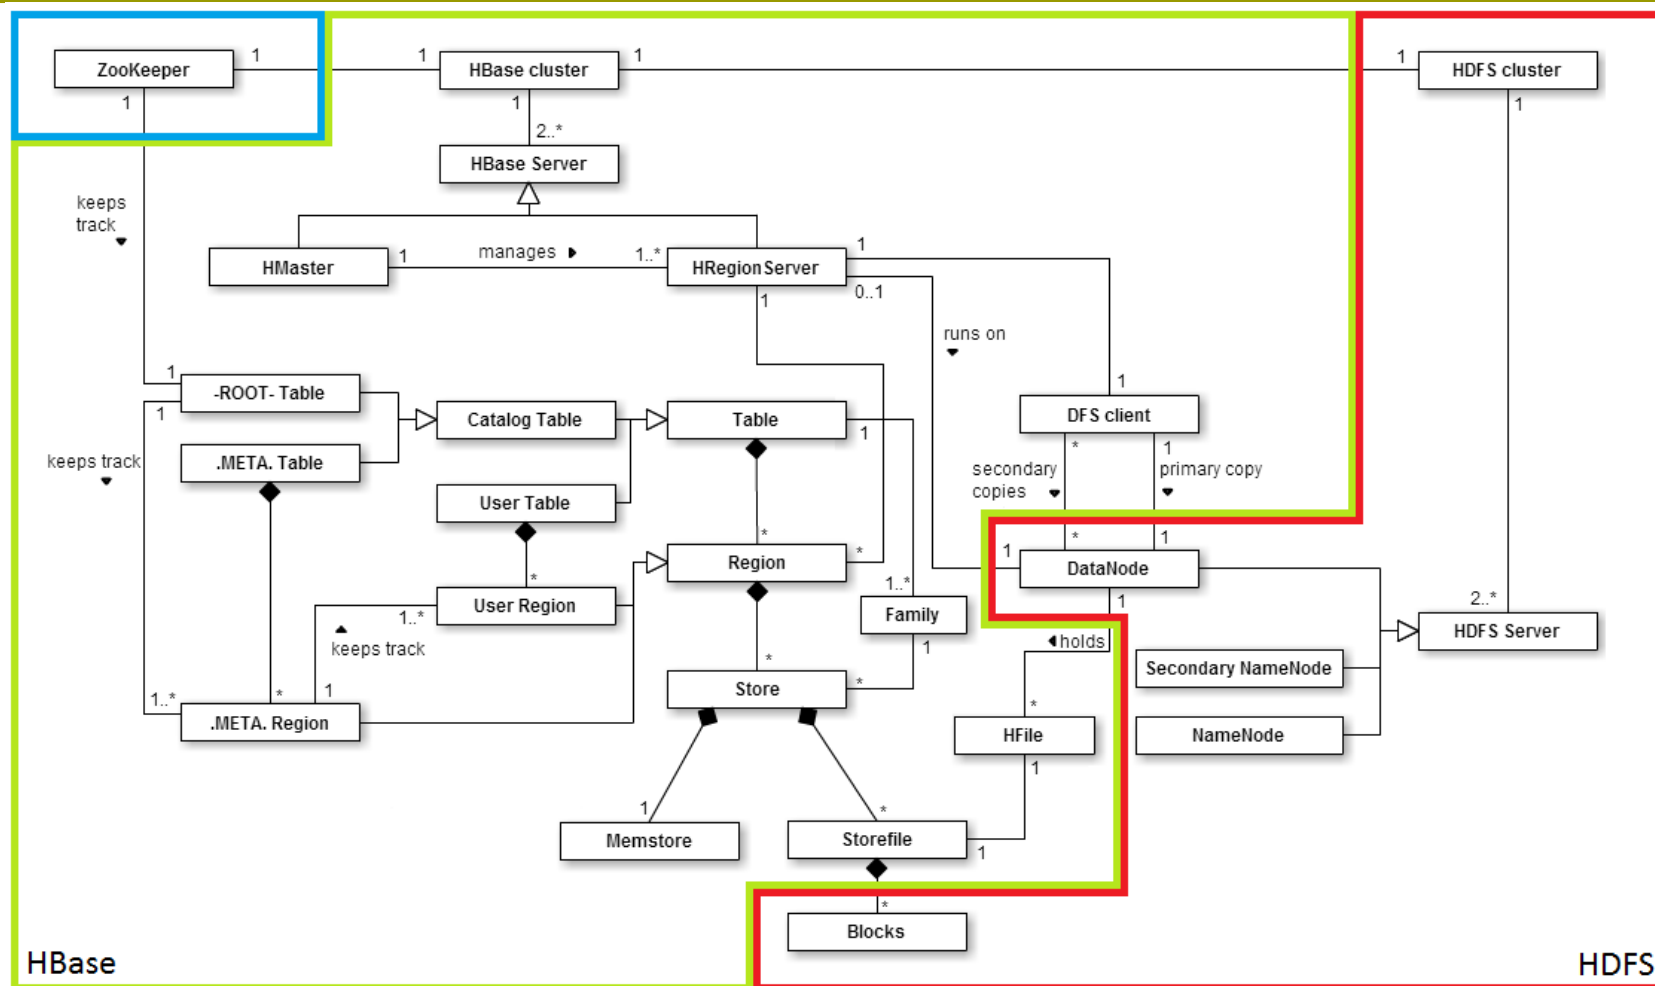


- Each table is horizontally fragmented into *Regions* (*tablets in BigTable*)
 - Dynamic fragmentation
 - By default into few hundreds of Mbs
 - Distributed on a cluster of machines or cloud
 - Uses the distributed B+ to decide in what region a tuple is stored
- At each Region rows are stored column-wise according to families (hybrid fragmentation)
 - In disk, as many files as families were defined
 - Static fragmentation (determined by the families; also determine data locality)
 - Block compression can be enabled (i.e., column families are compressed together)
- Massive usage of in-memory storage
 - Files in disk have an in-memory counterpart
 - Changes happen in main memory and are not flushed to disk until a compactation happens
 - A compactation merges the in-memory and disk files
- Metadata table (~ catalog)
 - It is the physical implementation of the catalog
 - Tuples are lexicographically sorted according to the key
 - Each row (entry) consists of <key, loc>
 - Key: it is the last key value in *that* Region
 - Loc: it is the physical address of a Region
 - This is a **distributed index cluster** (LSM-tree) on top of HDFS

Functional Components (I)

- ❑ Zookeeper
 - Quorum of servers that stores HBase system config info
- ❑ HMaster
 - Coordinates splitting of regions/rows across nodes
 - Controls distribution of HFile chunks
- ❑ Region Servers (HRegionServer)
 - Services HBase client requests
 - ❑ Manages stores containing all column families of the region
 - Logs changes
 - Guarantees “atomic” updates to one column family
 - Holds (caches) chunks of HFile into Memstores, waiting to be written
- ❑ HDFS
 - Stores all data including columns and logs
 - ❑ NameNode holds all metadata including namespace
 - ❑ DataNodes store chunks of a file
 - HBase uses two HDFS file types
 - ❑ HFile: regular data files (holds column data)
 - ❑ HLog: region’s log file (allows flush/fsync for small append-style writes)
- ❑ HFiles
 - Consist of large (e.g., 64MB) chunks
 - ❑ 3 copies of one chunk for availability (default)
- ❑ Clients
 - Read and write chunks
 - ❑ Locality & load determine which copy to access

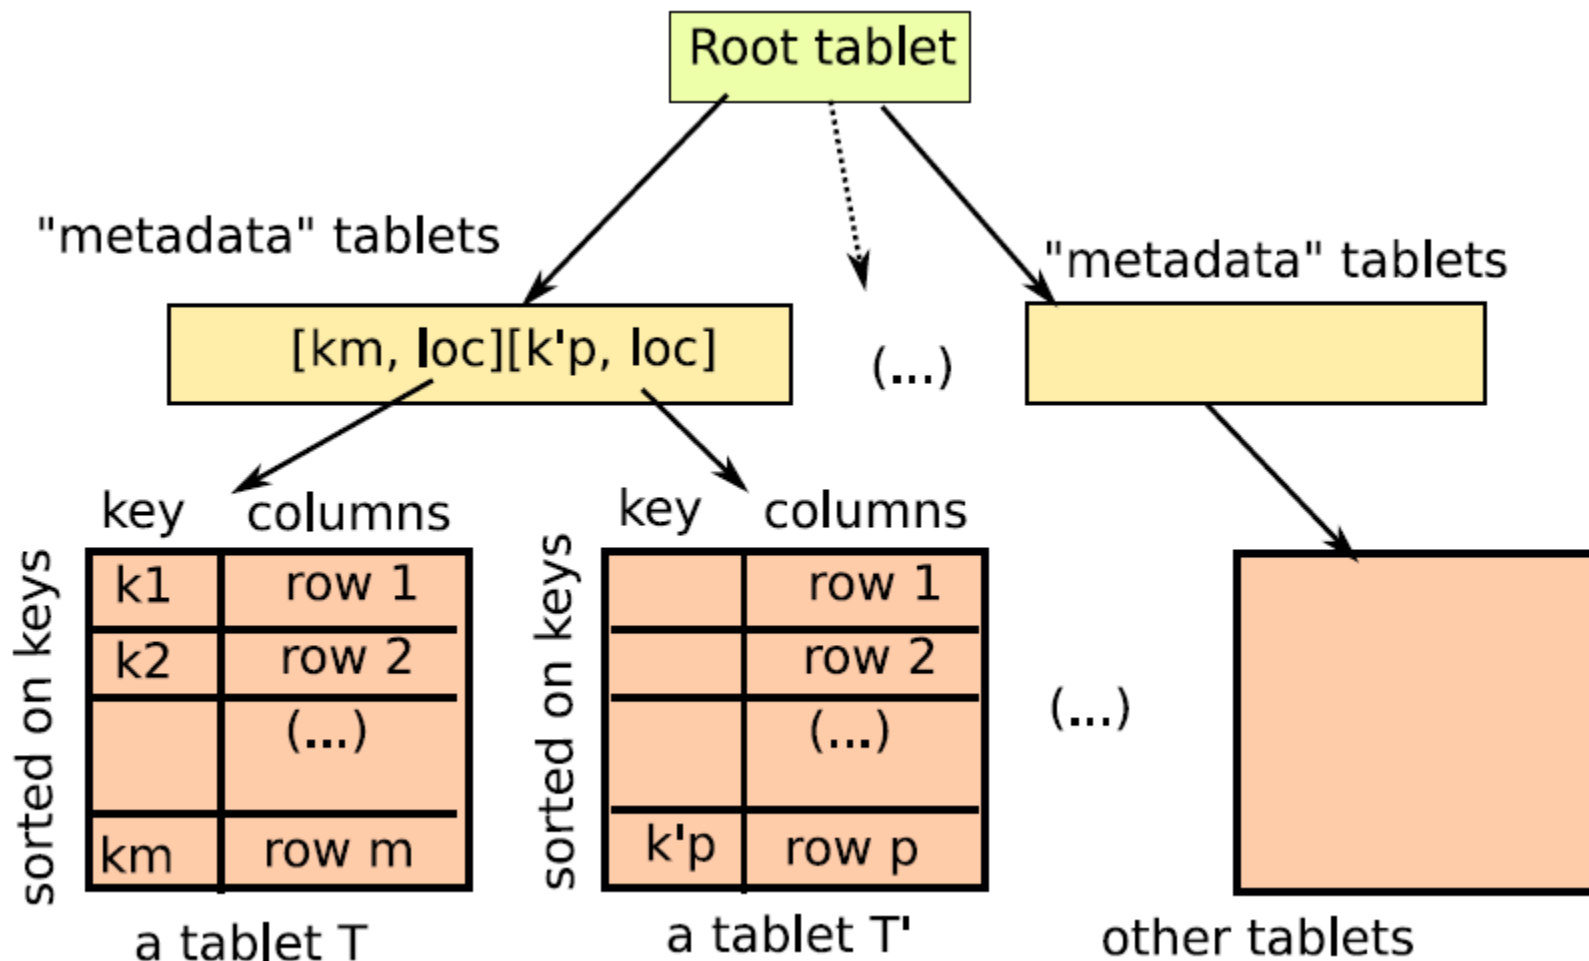
Functional Components (II)



- A primary copy must be stored in the same DataNode the HRegionServer runs on.
- Secondary copies can be stored in any DataNode different from the DataNode the HRegionServer runs on.
- All the stores of a given family correspond to the same table as this family.

$B \xrightarrow{n} \xrightarrow{m} A$ An element of class B is associated with n elements of A, and viceversa
 $B \longrightarrow \triangleright A$ Class B is a specialization (subtype) of class A
 $B \longrightarrow \blacklozenge A$ Class A is composed by elements of class B

HBase: A Distributed Index Cluster



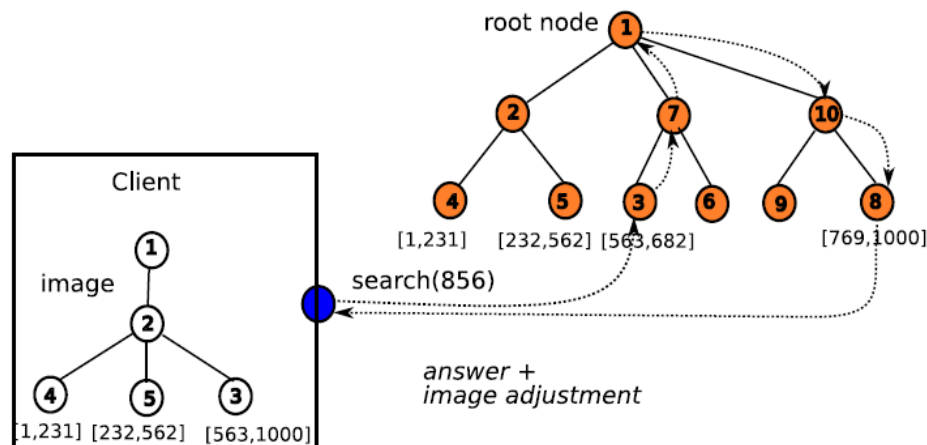
HBase Design Decisions (I)

- ❑ One master server
 - Maintenance of the table schemas
 - ❑ Root region
 - Monitoring of services (heartbeating)
 - Assignment of regions to servers
- ❑ Many region servers
 - Each handling around 100-1000 regions (i.e., horizontal fragments)
 - Apply concurrency and recovery techniques
 - **Managing split of regions**
 - ❑ A region server decides to split (e.g., >128MB)
 - ❑ Half of its regions are sent to another server
 - **Managing merge of regions**
- ❑ Client nodes

HBase Design Decisions (II)

- ❑ Split and merge affects the distributed tree, which must be updated
 - *Gossiping*
 - Lazy updates: discrepancies may cause out-of-range errors, which triggers a stabilization (**mistake compensation**) protocol
- ❑ Mistake compensation
 - The client keeps in cache the tree sent by the master and uses it to access data
 - If an out-of-range error is triggered, it is forwarded to the parent
 - ❑ In the worst case, 6 network round trips

HBase Design Decisions (II)



❑ Mistake compensation


- The client keeps in cache the tree sent by the master and uses it to access data
- If an out-of-range error is triggered, it is forwarded to the parent
 - ❑ In the worst case, 6 network round trips

Activity: Mistake Compensation


□ **Objective:** Understand how the global catalog is handled in HBase

□ **Tasks:**

1. (10') By pairs, solve the following exercise
 - a) Number of round trips if search(2602)
 - b) What is the expected number of round trips (i.e., in the average) for the next operation in the client
2. (-) Hand in the solution
3. (5') Discussion



Cache (@client)



Global catalog (@master)

HBase in a Nutshell

- ❑ HBase was thought to provide a specific answer for a specific problem
 - Architecture: Distributed index cluster
 - ❑ In-memory sorted map for new rows
 - ❑ Merge-sort for merging it with *old* data
 - Data structure: Table / row / families
 - ❑ Hybrid fragmentation (first horizontal, then vertical)
 - ❑ Thus, it cannot fully benefit from columnar processing
 - Concurrency: Timestamping MVCC (Multiversion)
 - ❑ Supports single-row transactions per file
 - Logging (RegionServers)
 - ❑ WAL (write-ahead protocol)
 - ❑ REDO login (No Force / No Steal)
 - CAP Theorem: CP

HBase Architecture

□ Refreshing the NOSQL Challenges

I. Distributed DB design

- Data fragments Horizontal fragmentation (fixed-size chunks)
- Data replication Replication performed by HDFS: Eager/ primary copy
- Node distribution Load balancing. Tuneable, by default depends on #RegionServers and size of the family file

II. Distributed DB catalog

- Fragmentation trade-off: Where to place the DB catalog
 - Global or local for each node
 - Centralized in a single node or distributed
 - Single-copy vs. Multi-copy

Global catalog: distributed tree
 Clustered data
 Centralized and multi-copy catalog
 (if several masters)
 Eager replication / secondary copy
 between them

III. Distributed query processing

- Data distribution / replication
- Communication overhead

MapReduce / Spark:

Data locality & parallelism
 Mostly query shipping but also data shipping
 Fault-tolerant

IV. Distributed transaction management

- How to enforce the ACID properties
 - Replication trade-off: Queries vs. Data consistency (updates)
 - Distributed recovery system
 - Distributed concurrency control system

CP

No transactions

HBase:

Concurrency: row level,
 multiversion timestamping
 Recovery: checkpointing logging

V. Security issues

- Network security Nothing!

HBase Architecture

□ NOSQL Goals

- Schemaless: No explicit schema [column-family key-value]
- Reliability / availability: Keep delivering service even if its software or hardware components fail [recovery] / [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth)

[distribution: CP]

Summary

- Goals of key-value stores
 - Schemaless
 - Consequences
 - Availability
 - Relationship to the CAP theorem
 - Scalability
 - By using the USL
 - Performance
 - Parallelize

Bibliography

- ❑ J. Dittrich et al. *Hadoop++: Making a yellow elephant run like a cheetah (without it even noticing)*. VLDB'10
- ❑ D. Jiang et al. *The performance of MapReduce: An In-depth Study*. VLDB'10
- ❑ Eric Brewer. *The CAP Theorem 12 years later*. <http://www.infoq.com/articles/cap-twelve-years-later-how-the-rules-have-changed>
- ❑ E. Brewer. *Towards Robust Distributed Systems*. Ann. ACM Symp. Principles of Distributed Computing (PODC 00). ACM, 2000
- ❑ F. Chang et al. *Bigtable: A Distributed Storage System for Structured Data*. OSDI'06
- ❑ Sanjay Ghemawat et al. *The Google File System*. SOSP'03
- ❑ Jeffrey Dean et al. *MapReduce: Simplified Data Processing on Large Clusters*. OSDI'04
- ❑ L. Liu and M.T. Özsu (Eds.). *Encyclopedia of Database Systems*. Springer, 2009
- ❑ N. J. Gunther. *A Simple Capacity Model of Massively Parallel Transaction Systems*. CMG National Conference, 1993