

Assume the following MapReduce program:

```
public void map(LongWritable key, Text value) {
    String line = value.toString();
    StringTokenizer tokenizer = new StringTokenizer(line);
    while (tokenizer.hasMoreTokens()) {
        write(new Text(tokenizer.nextToken()), new IntWritable(1));
    }
}

public void reduce(Text key, Iterable<IntWritable> values) {
    int sum = 0;
    for (IntWritable val : values) {
        sum += val.get();
    }
    write(key, new IntWritable(sum));
}
```

Consider the following data set:

- Block0: "a b b a c | c d c a e"
- Block1: "a b d d a | b b c c f"

Simulate the execution of the MapReduce code given the following configuration:

- The map and reduce functions are those of the wordcount
 - The combine function shares the implementation of the reduce
- One Split is one block
- The "|" divides the records inside each block
 - We have two records per block
- Hadoop is configured with the parameter *dfs.replication=1*
- We can keep four pairs [key,value] per spill
- We have two mappers and two reducers
 - Machine0, contains block0, runs mapper0 and reducer0
 - Machine1, contains block1, runs mapper1 and reducer1
- The hash function used to shuffle data to the reducers uses the correspondence:
 - {b,d,f}->0
 - {a,c,e}->1

Given Name: Family Name:

Given Name: Family Name:

Given Name: Family Name:

Fill the gaps in each step (numbers correspond to the phase in the MapReduce algorithm):

- 1) Machine0 contains blocks.
Machine1 contains blocks.
- 2) We keep replicas (including the master copy) per block.
- 3) We have splits per machine.
- 4) Mapper0 reads records.
Mapper1 reads records.

5) Spills in Machine0:

| Spill1 | Spill2 | Spill3 | Spill4 |
|------------|------------|------------|------------|
| [,][,] | [,][,] | [,][,] | [,][,] |
| [,][,] | [,][,] | [,][,] | [,][,] |

Spills in Machine1:

| Spill1 | Spill2 | Spill3 | Spill4 |
|------------|------------|------------|------------|
| [,][,] | [,][,] | [,][,] | [,][,] |
| [,][,] | [,][,] | [,][,] | [,][,] |

6) Partitions in machine0:

| Spill 1 | | Spill 2 | | Spill 3 | |
|-------------|-------------|-------------|-------------|-------------|-------------|
| Partition 0 | Partition 1 | Partition 0 | Partition 1 | Partition 0 | Partition 1 |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |

Partitions in Machine1:

| Spill 1 | | Spill 2 | | Spill 3 | |
|-------------|-------------|-------------|-------------|-------------|-------------|
| Partition 0 | Partition 1 | Partition 0 | Partition 1 | Partition 0 | Partition 1 |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |

7) Partitions in machine0:

| Spill 1 | | Spill 2 | | Spill 3 | |
|-------------|-------------|-------------|-------------|-------------|-------------|
| Partition 0 | Partition 1 | Partition 0 | Partition 1 | Partition 0 | Partition 1 |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |

Partitions in Machine1:

| Spill 1 | | Spill 2 | | Spill 3 | |
|-------------|-------------|-------------|-------------|-------------|-------------|
| Partition 0 | Partition 1 | Partition 0 | Partition 1 | Partition 0 | Partition 1 |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |
| [,][,] | [,][,] | [,][,] | [,][,] | [,][,] | [,][,] |

8) Files in machine0:

| File0 | File1 | File2 |
|----------------------|----------------------|----------------------|
| [,][,][,][,] | [,][,][,][,] | [,][,][,][,] |
| File3 | File4 | File5 |
| [,][,][,][,] | [,][,][,][,] | [,][,][,][,] |

Files in Machine1:

| File0 | File1 | File2 |
|----------------------|----------------------|----------------------|
| [,][,][,][,] | [,][,][,][,] | [,][,][,][,] |
| File3 | File4 | File 5 |
| [,][,][,][,] | [,][,][,][,] | [,][,][,][,] |

9) Merges in machine0:

| Merge0 | Merge1 |
|--------------------------------|--------------------------------|
| [,][,][,][,][,][,] | [,][,][,][,][,][,] |

Merges in Machine1:

| Merge0 | Merge1 |
|--------------------------------|--------------------------------|
| [,][,][,][,][,][,] | [,][,][,][,][,][,] |

10) Files in machine0:

| File0 | File1 |
|--------------------------------|--------------------------------|
| [,][,][,][,][,][,] | [,][,][,][,][,][,] |

Files in Machine1:

| File0 | File1 |
|--------------------------------|--------------------------------|
| [,][,][,][,][,][,] | [,][,][,][,][,][,] |

11) Reducer0 reads Files from machine0

and Files from machine1. (answer which Files)

Reducer1 reads Files from machine0

and Files from machine1. (answer which Files)

12) Merge in machine0:

| File0 | File1 |
|-----------------|-----------------|
| [,][,][,] | [,][,][,] |

Merge in Machine1:

| File0 | File1 |
|-----------------|-----------------|
| [,][,][,] | [,][,][,] |

13) Reduce function is executed times in machine0.

Reduce function is executed times in machine1.

14) Files in machine0:

| File0 | File1 |
|--------------------------------|--------------------------------|
| [,][,][,][,][,][,] | [,][,][,][,][,][,] |

Files in Machine1:

| File0 | File1 |
|--------------------------------|--------------------------------|
| [,][,][,][,][,][,] | [,][,][,][,][,][,] |